Lesson 6.2

Facilitating The AI-Native Value Workshop

AI Solution

AINOM
PROVIDED BY SCALED AGILE

Learning Objectives

By the end of this lesson, you will:

- Classify AI types (like sorting tools in a toolbox) and pick the right one for each job
- Define what your AI needs to do in production (what it actually does for real users)
- Design how humans and AI work together (like a pilot and autopilot sharing control)
- Evaluate your company's existing AI tools and show how reusing them saves time and money
- Help teams pick realistic AI approaches instead of chasing impossible dreams

WIIFM: Guide your team to the simplest, most realistic technical approach that delivers business value.

Section 2: AI Solution

Session 1 of 2: Foundation

Your New AI Teammate

5 min

Group Discussion: AI Joining Your Team

- An AI teammate joins your team next week. Which everyday task will you happily delegate—and how will your team use the time it wins back?

AI Solution Sub-Sections

- The AI Solution section of the blueprint connects business needs to the technical approach

Sub-section	Guidance
AI Classification	*Classify the AI application type (e.g., NLP, computer vision, predictive analytics), complexity level, and model selection approach.*
AI Capabilities	*Define essential AI functionalities required for production release to end users. These are the capabilities the system must have to function effectively in production. Define capability testing requirements and functional validation testing (Performance & Quality Benchmarks).*
Human-AI Interaction Design	*Design how humans and AI will collaborate. Define where humans add value and what requires human oversight. Consider usability testing, user acceptance testing, and human-AI collaboration testing (Performance & Quality Benchmarks).*
Performance & Quality	*Set quantified performance requirements for accuracy, speed, reliability, and quality thresholds. Define comprehensive testing strategies including unit tests, integration tests, performance tests, and user acceptance testing. Establish testing protocols for model validation, stress testing, and continuous monitoring in production.*

AI Classification - Four Ways to Build AI Solutions

Prompt-only

Leverages the model's vast pre-trained knowledge for general tasks. It is ideal for creative brainstorming and content generation where external data isn't needed.

RAG

Enterprise Sweet-Spot

Enriches the AI's response by providing it with relevant, real-time information retrieved from your own knowledge bases. This is the primary pattern for building context-aware enterprise solutions without costly model retraining.

Agentic

Orchestrates the AI to use multiple tools and perform a sequence of actions to achieve a complex goal. Best for automating end-to-end business processes like employee onboarding or complex customer support resolutions.

Fine-tuning

Adapts the model's fundamental parameters on a large, curated dataset to embed a highly specialized skill or style. This resource-intensive process is reserved for when RAG is insufficient and there is a critical need for unique model behavior at scale.

Guiding Principles

Simplicity First

Always verify if a simpler technology can satisfy the core business need before exploring more complex and costly options.

Time-to-Value is a Key Metric

Frame the decision in business terms. What is the value of an 85% solution delivered in one month versus a 95% solution delivered in six?

Govern Complexity

As your solution becomes more complex, the need for thorough testing, monitoring, and human-in-the-loop oversight grows significantly.

AI Capabilities - What Will Your AI Actually Do?

1. **Brainstorm the "Universe of Possibilities"**

   Walk through the user journey and identify every point where AI could assist.

   Common Capability Patterns: Summarization, Question Answering, Content Generation, Data Analysis, Task Automation.

2. **Filter Through the Lens of the Value Proposal**

   For every brainstormed capability, ask the critical question: "How does this feature help us achieve our primary business outcome?"

   If a capability doesn't clearly trace back to a key metric, it is not part of the MVP.

3. **Prioritize for a Minimum Viable Product (MVP)**

   Use a clear framework like **MoSCoW** to classify the filtered features:

   Must Have: Critical for the solution to be viable.

   Should Have: Important, but not required for initial launch.

   Could Have: A "nice to have" if time and resources permit.

   Won't Have: Explicitly out of scope for this release.

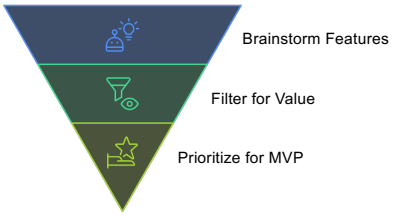

Human-AI Interaction Design - Defining the Partnership

	Researcher	Collaborator	Assistant	Agent
Level of Autonomy	AI presents information	AI is a "thought partner" for iterative work	AI handles delegated tasks	AI operates independently within set boundaries.
Role of AI	AI is a powerful tool for searching or research.	AI acts as a thought partner for brainstorming, analysis, and iterative refinement. The human and AI collaborate to shape the final output.	AI independently manages delegated tasks or provides completed work for human review. Automation begins here.	AI acts as a delegate, trusted to carry out multi-step processes.

Setting the "Rules of Engagement"

Confidence Threshold: At what level of certainty can the AI act alone?

Example: Below an 80% confidence score, the AI's output is flagged for mandatory human review.

Escalation Trigger: What specific situations must always involve a human?

Example: Any process involving sensitive personal data or transactions over a certain value must be handed off to a human employee.

Performance & Quality - Choosing the Right AI Engine

Model Category	Core Strengths	Key Consideration
State-of-the-Art Performance (e.g., OpenAI's GPT-5)	Complex reasoning, novel problem-solving, and applications requiring the absolute maximum raw capability.	Represents the highest performance tier, but also the highest cost and potential for unpredictable "emergent" behaviors.
High-Reliability & Safety-Focused (e.g., Anthropic's Claude 4, Google's Gemini Pro)	Enterprise-grade workflows, in-depth analysis of long documents, and applications where predictability and risk management are paramount.	Designed for more controllable and auditable outputs, making it a preferred choice for regulated or risk-averse environments.
Maximum Control & Customization (e.g., Meta's Llama 4, Mistral Large 2)	Cost-effective solutions at scale, deep model customization, and applications requiring full data privacy via on-premise or private cloud deployment.	Offers the most operational control and transparency, but requires significant internal technical expertise to manage, secure, and optimize.

Model Comparison: *LLM Leaderboard* (one of many such sites)

Additional Consideration - Making It Work at Scale

- **Cost (Token Usage):** Every piece of data sent to a model costs money.
  - How can we design our solution to be efficient and affordable?
  - Example:1 Million tokens/day × $0.02/1k tokens = $20/day ($7,200/year)
- **Speed (Latency):** How quickly does the user need an answer?
  - Does this require an instant response, or can it take time?
  - Example: A real-time chatbot vs. an overnight data analysis report.
- **Quality (Accuracy):** What is the business cost of the AI being wrong?
  - What level of precision is truly necessary to deliver business value?
  - Example: 85% for drafting marketing ideas vs. 99.9% for medical diagnostics.
- **Volume (Scale):** How many users and requests will the system handle over time?
  - Are we building for Day 1 usage or for Year 2 growth?
  - Example: A pilot for 100 internal users vs. a system for 10,000 customers.

Ready for Next Steps

We've:

- **Explored** classification approaches and trade-offs
- **Defined** core capabilities mapped to business value
- **Understood** human-AI interaction patterns and boundaries

Next session (after a practical exercise):

- How to facilitate these technical discussions

6.2 What's the Pattern?

15 min

Companion App

Group Activity: What's the Pattern?

- Part 1: Individual Classification (5 minutes)
  - Read through the 25 stakeholder statements. Pick up to 10 that resonate with you.
  - For each statement, decide which of the five AI patterns is the "best fit" to solve the underlying problem:

1. Traditional ML (Forecasting, Pattern Recognition, Computer Vision)
2. Prompt-only (chatbot)
3. RAG (Retrieval-Augmented Generation – specific context)
4. Fine-Tuning (domain-specific model modification)
5. Agentic (workflows, tools, and automation)

  - Drag and drop each statement into the corresponding category.
- Part 2: Group Discussion & Consensus (10 minutes)
  - In your group, compare your individual classifications.
  - Discuss any disagreements. Your goal is to reach a group consensus on the best-fit pattern for all 25 statements.
  - For each pattern, be ready to explain the "why" behind your choice. What keywords or phrases in the statement led you to your conclusion?

Connect the practice to the strategy

5 min

Class Discussion: Finding the Best Fit

- During this exercise, you weren't just sorting statements; you were training your ear. You learned to spot keywords like "based only on our documents," "learn our unique style," or "automatically do X, then Y" that point to a specific technical path.
- **Question:** Let's discuss as a class: As a Change Agent facilitating a workshop, how does having this 'trained ear' allow you to guide the conversation more effectively and keep it from going off track?

Section 2: AI Solution

Session 2 of 2: Facilitation

From Demo to Disaster

5 min

Group Discussion: The Production Reality Check

- Imagine a data scientist shows you a brilliant demo of an AI on their laptop. It uses an open-source model to perfectly summarize long, complex legal documents. Everyone is amazed. Your business sponsor wants to roll it out to all 300 lawyers in the company next month.
- In your group, brainstorm: What technical realities and hidden costs might turn this 'brilliant demo' into a production nightmare?

The AI-Native Change Agent as Facilitator

- **Shift focus:**

  Understanding the tech → Guiding the conversation

- **Your role:**

  Extract decisions, not become the architect

- **Key principle:**

  Good questions > Perfect answers

- **Remember:**

  You're facilitating experts who know more than you

What Works (and What Doesn't)

Patterns for Success

1. **Target a narrow, high-value problem**
   Solve a specific, well-defined business pain point first. A narrow victory is far better than a broad defeat.
2. **Use the simplest tool for the job**
   Start with the most straightforward technical approach that meets the need. Complexity is not a measure of success.
3. **Build a feedback loop from day one**
   Design clear mechanisms for users to report issues and for you to measure the AI's performance. The goal is to learn and improve rapidly.

Anti-Patterns to Avoid

1. **Don’t try to “boil the ocean”**
   Avoid overly ambitious "Phase 1" projects that try to solve everything at once. This is the fastest path to the "Proof-of-Concept Graveyard."
2. **Don’t use a sledgehammer to crack a nut**
   If a problem can be solved effectively with simple business rules or traditional software, do not force an AI solution.
3. **Don’t neglect the human in the loop**
   Never assume 100% AI accuracy. Always design for human oversight, clear escalation paths, and the ability for a person to intervene.

AI Classification - Guiding the "Which Type?" Conversation

- **Blueprint asks:**
  "What type of AI are we building?"
- **Effective answer:**
  "RAG system for over 10K product docs to answer customer questions in <30 seconds"
- **Common Pitfall:**
  "We need AI" or "Let's use ChatGPT for everything."
- **Your questions unlock clarity**

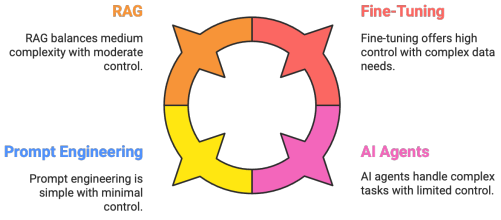

AI Capabilities - From Wishes to Features

- **Blueprint asks:**

  "What must this AI actually do?"

- **Effective answer:**

  1. "Answer FAQs about product specifications (Must Have),
  2. Generate a draft of the weekly performance report (Nice to Have), and
  3. Route complex support tickets to the correct expert team (Must Have)."

- **Common Pitfall:**

  "Make it smart" or listing 20 features for MVP

- Use MoSCoW to prioritize Features, and combine with methods like KANO to remain customer-focused

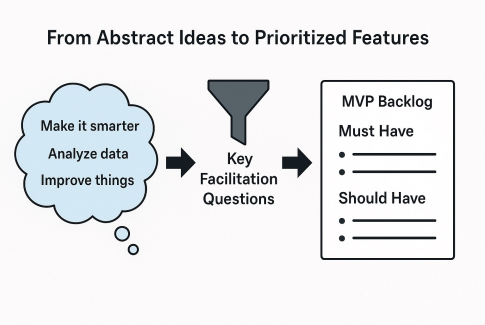

Human-AI Interaction - The Handoff Dance

- **Blueprint asks:**

  "When does AI hand off to humans?"

- **Effective answer:**

  “The AI must flag missing standard NDAs in 95% of contracts. False positives under 10% are acceptable but any >15% should trigger a manual audit.”

  → Sets clear accuracy and false-positive/error thresholds, with monitoring triggers—great for responsible deployment.

- **Common Pitfall:**

  “It should catch all missing clauses perfectly.”

  → Unrealistic and untestable; doesn’t guide build vs. buy or evaluation criteria.

- Find the sweet spot for your context

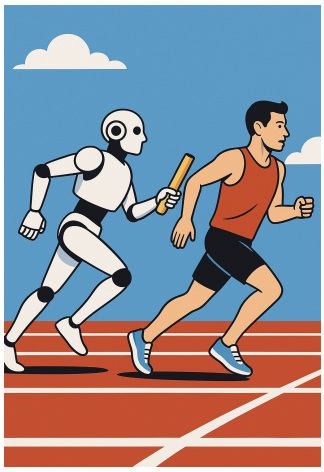

Performance and Quality - Defining "Good Enough"

- **Blueprint asks:**

  "What performance is acceptable?"

- **Effective answer:**

  "95% accuracy on top 100 questions, <2 second response, handles 1000 users"

- **Common Pitfall:**

  "Fast and accurate" or "100% accuracy always"

- Trade-offs make or break your project

Perfection delays takeoff!

Additional Considerations: Sharp Scope & Visible Risks

- **Watch for Drift**
  "While we're at it…" is the enemy of progress. Anchor ideas to success metrics and the 12-week timeline.

- **Ask These Early and Often**

  "Does this directly contribute to our value goal?"

  "What could go wrong?"

  "What do we absolutely need to learn before we commit?"

- **Make Confidence & Trade-offs Explicit**
  - *"Fist of five—how confident are we in this plan?"*
  - *"If we add this, what comes out?"*
  - Assign High/Med/Low confidence to major decisions

- **Write Down, Don't Work Around**
  Log risks and new ideas. Don't solve or chase—park and prioritize.

- **Use the Three-Column Discipline**
  - **MVP** — what we commit to now
  - **Next Phase** — near-term possibilities
  - **Future Vision** — capture the ambition, don't derail the plan

Ready to Build

- **Summarize:**

  Classification approach, core capabilities, human-ai interaction, and performance & quality

- **Confirm:**

  "Do we have enough to move forward , or are there showstoppers?"

- **Note dependencies:**

  "What do we need from Data/Ops/Risk sections?"

- **Celebrate progress:**

  "We've designed something that could actually work!"

6.3 AI Solution – Facilitation Challenge

15 min

Companion App

Group Activity:

- **Facilitation Principle:** A Change Agent's AI fluency helps guide teams to the most achievable path, not just the most technically impressive one.
- **Assign Your Roles:** Your group designates a Lead Facilitator, Scribe, and Clarifier.
- **In the Companion App, you will**
  - **Analyze the following:**
    - The Scenario
    - The team's flawed sub-section
    - Voices from the Meeting
  - **Brainstorm** powerful questions based on your analysis.
  - **Record** your top questions and **explain** the 'why' behind them.

Translation Excellence

5 min

Class Discussion: AI Solution Debrief

- Our Objective Was to highlight that skillful facilitation translates complex technical choices into clear business trade-offs.
- Let's share the strategic thinking behind the questions you prepared. We'll hear from a couple of groups, and I'm particularly interested in any different questions or approaches you took.

Lesson Review

You can facilitate the AI Solution section of the AI-Native Value Blueprint:

- Classify AI types and pick the right one for each job
- Define what your AI needs to do in production
- Design how humans and AI work together effectively
- Evaluate your company's existing AI tools to maximize value
- Help teams pick realistic AI approaches that actually work

Insights & Action

3 min

- Reflect

  Think about a recent meeting where technical and business teams struggled to understand each other. Write about how having AI fluency might have changed that conversation. What questions could you have asked to bridge the gap?

- Apply

  Ask AI: "I need to explain [insert AI concept like RAG or fine-tuning] to my non-technical manager. You're my communication coach. Help me create a simple analogy and a 2-minute explanation that avoids jargon while keeping the key points accurate."

- Explore

  LLM Leaderboard — Compare different AI models based on performance parameters
  What Have We Learned about Human-AI Collaboration? – BCG
  Seizing the Agentic AI Advantage – McKinsey
  AI agents: Opportunities, risks, and mitigations – IBM Ethics Board